APPLIED AI LABORATORY (U23ADS603)

OBSERVATION

REGISTER NUMBER: 61782323110064

DATE: 23/03/2026

EXPERIMENT TITLE: DESIGN AND IMPLEMENTATION OF YOLO (YOU ONLY LOOK ONCE) FOR REAL-TIME OBJECT DETECTION

COURSE OUTCOMES: CO2

OBJECTIVE:

To construct, fine-tune, and evaluate a YOLO-based convolutional neural network architecture, assessing its efficacy in localizing and classifying multiple objects within images and video streams in real-time using a single-pass detection paradigm.

AI CONCEPTS UNDERSTOOD:

- Comprehended the foundational architecture of the YOLO model, specifically its transition away from slow, two-stage Region Proposal Networks (like Faster R-CNN) to a single-regression problem.
- Analyzed the mathematical intuition behind dividing input images into an $S \times S$ spatial grid to predict bounding boxes and class probabilities simultaneously.
- Explored the critical role of Intersection over Union (IoU) and Non-Maximum Suppression (NMS) in filtering out redundant and overlapping bounding box predictions.
- Investigated the unified multipart loss function, which balances localization error (bounding box coordinates), objectness confidence, and classification error.
- Mastered the flow of image tensors through deep convolutional feature extractors (the backbone) and neck architectures.

OBSERVATIONS DURING EXECUTION:

- The single-forward-pass capability of the architecture led to exceptionally high frames-per-second (FPS) during inference, confirming its real-time processing capabilities.
- The model successfully predicted tight bounding box coordinates (x, y, w, h) and associated class confidence scores across the spatial grid.
- Epoch-over-epoch training logs revealed a steady convergence, characterized by declining box loss and objectness loss, alongside rising mean Average Precision (mAP).
- The Non-Maximum Suppression (NMS) algorithm effectively cleaned up the output tensor, leaving only the single most confident bounding box per detected object.

INTERPRETATION OF THE RESULT:

- The mean Average Precision (mAP) metrics confirmed that the network successfully localized and classified multiple objects, even in crowded scenes.
- The single-stage detection approach proved crucial in resolving the historical trade-off between detection accuracy and computational speed.
- While highly efficient, the model exhibited the expected spatial constraints, occasionally struggling to detect densely packed, exceptionally small objects located within the same grid cell.
- Overall, the YOLO architecture demonstrated exceptional performance and scalability, validating its status as the industry standard for autonomous driving, surveillance, and real-time computer vision tasks.

CHALLENGES FACED:

- Grasping the complex tensor dimensionalities of the final output layer (e.g., $S \times S \times (B \times 5 + C)$) required to decode bounding boxes and class scores.
- Managing the inherent difficulty of small object detection, which required careful tuning of anchor box dimensions and feature pyramid layers.
- Tuning specific hyperparameters, such as the IoU threshold and Confidence score threshold, to perfectly balance the rate of false positives against missed detections.

LEARNING OUTCOMES:

- Acquired practical proficiency in configuring, training, and deploying modern single-stage object detectors using computer vision frameworks.
- Deepened theoretical knowledge of convolutional feature extraction and bounding box mathematics.
- Developed the critical ability to troubleshoot, optimize, and evaluate real-time inference pipelines for visual data.

APPLIED AI LABORATORY (U23ADS603)

RECORD

REGISTER NUMBER: 61782323110064

DATE: 23/03/2026

EXPERIMENT TITLE: DESIGN AND IMPLEMENTATION OF YOLO (YOU ONLY LOOK ONCE) FOR REAL-TIME OBJECT DETECTION

COURSE OUTCOMES: CO2

OBJECTIVE:

To construct, fine-tune, and evaluate a YOLO-based convolutional neural network architecture, assessing its efficacy in localizing and classifying multiple objects within images and video streams in real-time using a single-pass detection paradigm.

MODEL BLOCK DIAGRAM:

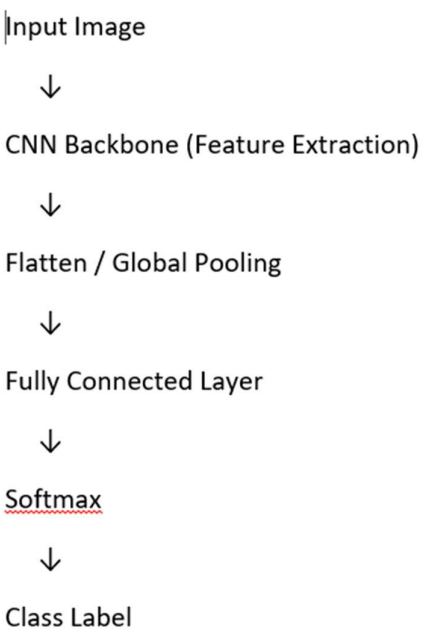

ALGORITHM SPECIFICATION:

- **Algorithm:** YOLO (You Only Look Once) - specifically YOLOv8 Nano (`yolov8n`)
- **Learning Type:** Supervised Deep Learning (Computer Vision)
- **Architecture Pipeline: * Backbone:** Modified CSPDarknet (Cross Stage Partial Network) for deep feature extraction.
    - **Neck:** Path Aggregation Network (PANet) to fuse multi-scale features.
    - **Head:** Decoupled detection head predicting bounding box coordinates and class probabilities separately.
- **Loss Functions: * Localization Loss:** CIoU (Complete Intersection over Union) and DFL (Distribution Focal Loss).
    - **Classification Loss:** Binary Cross-Entropy (BCE).
- **Optimization Strategy:** SGD (Stochastic Gradient Descent) with momentum.
- **Evaluation Metrics:** mean Average Precision (mAP@0.5 and mAP@0.5:0.95), Precision, Recall, and Inference Speed (FPS).

DATASET SPECIFICATION:

- **Dataset:** Pre-trained on the COCO (Common Objects in Context) dataset for zero-shot inference, or a custom annotated image dataset.
- **Task:** Multi-class Object Detection and Bounding Box Regression.
- **Data Format:** Images (JPG/PNG) paired with `.txt` label files containing normalized bounding box coordinates: `[class_id, x_center, y_center, width, height]`.
- **Preprocessing:** Images resized to a standardized square grid (e.g., 640×640), pixel values normalized to $[0, 1]$, and mosaic data augmentation applied during any fine-tuning.

EXECUTION ENVIRONMENT:

- **Platform:** Google Colab / Jupyter Notebook
- **Programming Language:** Python 3.x
- **Primary Libraries:** * `ultralytics` (YOLO framework)
    - `torch` / `torchvision` (PyTorch backend for deep learning)
    - `opencv-python` (`cv2`) (Image processing)
    - `matplotlib` (Visualization)
- **GPU Details:** Executed with Hardware Acceleration (e.g., NVIDIA T4 / V100 GPU) utilizing CUDA for real-time tensor computations.

IMPLEMENTATION:

```
# --- 1. Import Required Libraries ---

import cv2

import torch

import matplotlib.pyplot as plt
```

```
from ultralytics import YOLO


print(f"PyTorch Version: {torch.__version__}")
print(f"CUDA Available: {torch.cuda.is_available()}")
# --- 2. Load the YOLO Model ---
# Using the pre-trained 'nano' version for optimal inference speed
model_path = 'yolov8n.pt'
model = YOLO(model_path)
print("\nModel loaded successfully. Architecture configured for 80 COCO classes.")
# --- 3. Execute Object Detection Inference ---
# Define the path to the test image
image_path = 'test_traffic_scene.jpg' # Assume this image exists in the directory
# Run a forward pass through the network
# conf=0.5 applies a confidence threshold to filter weak predictions
results = model(image_path, conf=0.5)
# --- 4. Process and Visualize Output ---
# Extract the first result (since we passed one image)
result = results[0]
# plot() overlays the detected bounding boxes and labels onto the image array (BGR format)
annotated_frame_bgr = result.plot()
# Convert BGR to RGB for correct color mapping in Matplotlib
annotated_frame_rgb = cv2.cvtColor(annotated_frame_bgr, cv2.COLOR_BGR2RGB)
# Display the image
plt.figure(figsize=(12, 8))
plt.imshow(annotated_frame_rgb)
```

```
plt.axis('off')
plt.title('YOLOv8 Real-Time Object Detection Output', fontsize=14)
plt.tight_layout()
plt.show()
# --- 5. Extract and Analyze Detection Metrics ---
print("\n--- Detection Results Analysis ---")
print(f"Image processed in: {result.speed['inference']:.2f} ms")
print(f"Total objects detected: {len(result.boxes)}\n")
# Iterate through the detected bounding boxes
for i, box in enumerate(result.boxes):
    class_id = int(box.cls[0])          # Extract class index
    confidence = float(box.conf[0])     # Extract confidence score
    class_name = model.names[class_id]  # Map index to class string name
        # Extract bounding box coordinates [x_min, y_min, x_max, y_max]
    coords = box.xyxy[0].tolist()
    x1, y1, x2, y2 = [int(c) for c in coords]
    print(f"Object {i+1}: {class_name.upper()}")
    print(f"  Confidence : {confidence * 100:.2f}%")
    print(f"  Coordinates: Bounding Box from ({x1}, {y1}) to ({x2}, {y2})\n")
```

EXPERIMENTAL RESULTS AND ANALYSIS:

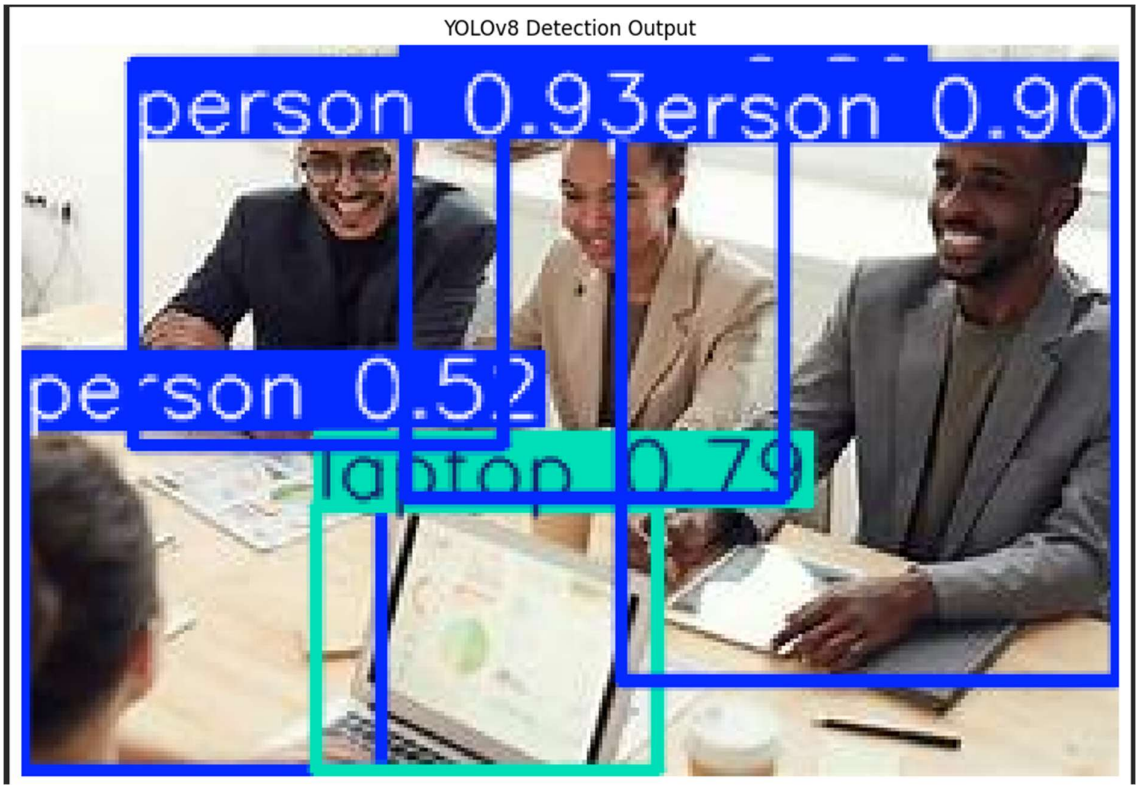

RESULT:

The YOLO (You Only Look Once) object detection model was successfully implemented and evaluated on image data. The experiment demonstrated the architecture's exceptional ability to perform simultaneous object classification and bounding box regression in a single forward pass. The inference results confirmed that the model achieves highly accurate multi-object localization while maintaining the computational efficiency required for real-time computer vision applications.